

Choosing a college to attend is one of the most important decision one can make. What measures the success of a college degree? It can be either the height of the societal ladder that you can climb, or how much content you are with your life. For the sake of this project, we are measuring a quantitative variable: Salary.

The college in exploration are divided based on 3 categories:

- Colleges by Type
- Colleges by Region
- Salary by major

The dataset is from The Wall Street Journal, available on Kaggle at: <https://www.kaggle.com/wsj/college-salaries> (<https://www.kaggle.com/wsj/college-salaries>).

First, let's import some library we will use later

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Let's load our first dataset: Colleges by Type

```
In [2]: college_type = pd.read_csv("salaries-by-college-type.csv")
college_region = pd.read_csv("salaries-by-region.csv")
```

Take a peek at our data

```
In [3]: college_type.head(5)
```

Out[3]:

	School Name	School Type	Starting Median Salary	Mid-Career Median Salary	Mid-Career 10th Percentile Salary	Mid-Career 25th Percentile Salary	Mid-Career 75th Percentile Salary
0	Massachusetts Institute of Technology (MIT)	Engineering	\$72,200.00	\$126,000.00	\$76,800.00	\$99,200.00	\$168,000.00
1	California Institute of Technology (CIT)	Engineering	\$75,500.00	\$123,000.00	NaN	\$104,000.00	\$161,000.00
2	Harvey Mudd College	Engineering	\$71,800.00	\$122,000.00	NaN	\$96,000.00	\$180,000.00
3	Polytechnic University of New York, Brooklyn	Engineering	\$62,400.00	\$114,000.00	\$66,800.00	\$94,300.00	\$143,000.00
4	Cooper Union	Engineering	\$62,200.00	\$114,000.00	NaN	\$80,200.00	\$142,000.00

In [4]: `college_region.head(5)`

Out[4]:

	School Name	Region	Starting Median Salary	Mid-Career Median Salary	Mid-Career 10th Percentile Salary	Mid-Career 25th Percentile Salary	Mid-Career 75th Percentile Salary	Number of Applicants
0	Stanford University	California	\$70,400.00	\$129,000.00	\$68,400.00	\$93,100.00	\$184,000.00	\$2
1	California Institute of Technology (CIT)	California	\$75,500.00	\$123,000.00	NaN	\$104,000.00	\$161,000.00	NaN
2	Harvey Mudd College	California	\$71,800.00	\$122,000.00	NaN	\$96,000.00	\$180,000.00	NaN
3	University of California, Berkeley	California	\$59,900.00	\$112,000.00	\$59,500.00	\$81,000.00	\$149,000.00	\$2
4	Occidental College	California	\$51,900.00	\$105,000.00	NaN	\$54,800.00	\$157,000.00	NaN

Notice that the salary column is identical, let's combine these 2 dataset so we can have both location and school type for each college

In [5]: `df = pd.merge(college_type, college_region)`

In [6]: `df.head()`

Out[6]:

	School Name	School Type	Starting Median Salary	Mid-Career Median Salary	Mid-Career 10th Percentile Salary	Mid-Career 25th Percentile Salary	Mid-Career 75th Percentile Salary
0	Massachusetts Institute of Technology (MIT)	Engineering	\$72,200.00	\$126,000.00	\$76,800.00	\$99,200.00	\$168,000.00
1	California Institute of Technology (CIT)	Engineering	\$75,500.00	\$123,000.00	NaN	\$104,000.00	\$161,000.00
2	Harvey Mudd College	Engineering	\$71,800.00	\$122,000.00	NaN	\$96,000.00	\$180,000.00
3	Polytechnic University of New York, Brooklyn	Engineering	\$62,400.00	\$114,000.00	\$66,800.00	\$94,300.00	\$143,000.00
4	Cooper Union	Engineering	\$62,200.00	\$114,000.00	NaN	\$80,200.00	\$142,000.00

Ah, the region is there but it's at the last column, which is not visible much. Let's move it to the second column

In [7]: `region = df.pop("Region")  
df.insert(1, 'Region', region)`

In [8]: `df.head()`

Out[8]:

	School Name	Region	School Type	Starting Median Salary	Mid-Career Median Salary	Mid-Career 10th Percentile Salary	Mid-Career 25th Percentile Salary
0	Massachusetts Institute of Technology (MIT)	Northeastern	Engineering	\$72,200.00	\$126,000.00	\$76,800.00	\$99,200.00
1	California Institute of Technology (CIT)	California	Engineering	\$75,500.00	\$123,000.00	NaN	\$104,000.00
2	Harvey Mudd College	California	Engineering	\$71,800.00	\$122,000.00	NaN	\$96,000.00
3	Polytechnic University of New York, Brooklyn	Northeastern	Engineering	\$62,400.00	\$114,000.00	\$66,800.00	\$94,300.00
4	Cooper Union	Northeastern	Engineering	\$62,200.00	\$114,000.00	NaN	\$80,200.00

Nice! Now let's take a quick look overview of our data

In [9]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 260 entries, 0 to 259
Data columns (total 9 columns):
School Name                260 non-null object
Region                    260 non-null object
School Type                260 non-null object
Starting Median Salary     260 non-null object
Mid-Career Median Salary   260 non-null object
Mid-Career 10th Percentile Salary  223 non-null object
Mid-Career 25th Percentile Salary  260 non-null object
Mid-Career 75th Percentile Salary  260 non-null object
Mid-Career 90th Percentile Salary  223 non-null object
dtypes: object(9)
memory usage: 20.3+ KB
```

In [10]: `df.describe()`

Out[10]:

	School Name	Region	School Type	Starting Median Salary	Mid-Career Median Salary	Mid-Career 10th Percentile Salary	Mid-Career 25th Percentile Salary
<b>count</b>	260	260	260	260	260	223	260
<b>unique</b>	240	5	5	142	162	135	172
<b>top</b>	Pennsylvania State University (PSU)	Northeastern	State	\$42,600.00	\$72,100.00	\$40,100.00	\$54,100.00
<b>freq</b>	2	69	169	7	5	6	6

Notice that all the variables are object type. Let's change School Name, Region, School Type to string, and the rest to numeric values

Before we change it to numeric values, we need to drop the "\$" sign before the value

In [11]: `# Remove the dollar $ sign  
colstocheck = df.columns  
df[colstocheck] = df[colstocheck].replace({'\': '', ',': ''}, regex = True)`

In [17]: `to_change = ["Starting Median Salary", "Mid-Career Median Salary", "Mid-Career 10th Percentile Salary", "Mid-Career 25th Percentile Salary", "Mid-Career 75th Percentile Salary", "Mid-Career 90th Percentile Salary"]  
for i in to_change:  
 df[i] = df[i].astype(float)  
  
#df["Mid-Career 90th Percentile Salary"] = df["Mid-Career 90th Percentile Salary"].astype(float)`

In [18]: `df.head()`

Out[18]:

	School Name	Region	School Type	Starting Median Salary	Mid-Career Median Salary	Mid-Career 10th Percentile Salary	Mid-Career 25th Percentile Salary	Per
0	Massachusetts Institute of Technology (MIT)	Northeastern	Engineering	72200.0	126000.0	76800.0	99200.0	168
1	California Institute of Technology (CIT)	California	Engineering	75500.0	123000.0	NaN	104000.0	161
2	Harvey Mudd College	California	Engineering	71800.0	122000.0	NaN	96000.0	180
3	Polytechnic University of New York Brooklyn	Northeastern	Engineering	62400.0	114000.0	66800.0	94300.0	143
4	Cooper Union	Northeastern	Engineering	62200.0	114000.0	NaN	80200.0	142

In [19]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 260 entries, 0 to 259
Data columns (total 9 columns):
School Name                260 non-null object
Region                    260 non-null object
School Type                260 non-null object
Starting Median Salary      260 non-null float64
Mid-Career Median Salary    260 non-null float64
Mid-Career 10th Percentile Salary  223 non-null float64
Mid-Career 25th Percentile Salary  260 non-null float64
Mid-Career 75th Percentile Salary  260 non-null float64
Mid-Career 90th Percentile Salary  223 non-null float64
dtypes: float64(6), object(3)
memory usage: 20.3+ KB
```

Great! Looks like we have pretty much done with the pre-processing data. Now the last step would be to process the missing values.

There are several ways to handle missing data:

1. Delete the entire rows which have the missing values. This is the simplest, but not ideal because the more data we have for this project, the better! There are 15% of the data go missing in this dataset.
2. Fill in the missing values with the mean, median, or mode of the same type of variable. Let's go with this option

Actually after doing some research, there are 6 ways to handle this! Let's try some methods recommended here: <https://towardsdatascience.com/6-different-ways-to-compensate-for-missing-values-data-imputation-with-examples-6022d9ca0779> (<https://towardsdatascience.com/6-different-ways-to-compensate-for-missing-values-data-imputation-with-examples-6022d9ca0779>)

```
In [20]: #import sys
          #from impute.imputation.cs import fast_knn
          #sys.setrecursionlimit(100000) #Increase the recursion limit of the OS

          # start the KNN training
          #imputed_training=fast_knn(train.values, k=30)

-----
ModuleNotFoundError                                Traceback (most recent call last)
<ipython-input-20-79d38c8586f7> in <module>()
      1 import sys
----> 2 from impute.imputation.cs import fast_knn
      3 sys.setrecursionlimit(100000) #Increase the recursion limit of the OS
      4
      5 # start the KNN training

ModuleNotFoundError: No module named 'impute'
```



```
In [27]: # Method 1: Replacing NaN values with mean value of that variable
#Impute the values using scikit-learn SimpleImpute Class
#####import sklearn
#####from sklearn.preprocessing import Imputer
#####imputer = Imputer(missing_values="NaN", strategy="mean", axis=1)
#####imputed_mean_df = imputer.fit_transform(df)
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-27-4ecc358f4152> in <module>()
      4 from sklearn.preprocessing import Imputer
      5 imputer = Imputer(missing_values="NaN", strategy="mean", axis=1)
----> 6 imputed_mean_df = imputer.fit_transform(df)
      7 #from sklearn.impute import SimpleImputer
      8 #imp_mean = SimpleImputer(strategy='mean') #for median imputation replace 'mean' with 'median'

~\Anaconda3\lib\site-packages\sklearn\base.py in fit_transform(self, X, y, **fit_params)
    515         if y is None:
    516             # fit method of arity 1 (unsupervised transformation)
--> 517             return self.fit(X, **fit_params).transform(X)
    518         else:
    519             # fit method of arity 2 (supervised transformation)

~\Anaconda3\lib\site-packages\sklearn\preprocessing\imputation.py in transform(self, X)
    321         else:
    322             X = check_array(X, accept_sparse='csr', dtype=FLOAT_DTYPES,
S,
--> 323                             force_all_finite=False, copy=self.copy)
    324
    325             if sparse.issparse(X):

~\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_array(array, accept_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, ensure_min_samples, ensure_min_features, warn_on_dtype, estimator)
    431                 force_all_finite)
    432         else:
--> 433             array = np.array(array, dtype=dtype, order=order, copy=copy)
    434
    435             if ensure_2d:
```

**ValueError:** could not convert string to float: 'State'

```
In [31]: df.isnull().sum()
```

```
Out[31]: School Name      0
         Region          0
         School Type      0
         Starting Median Salary  0
         Mid-Career Median Salary  0
         Mid-Career 10th Percentile Salary  37
         Mid-Career 25th Percentile Salary  0
         Mid-Career 75th Percentile Salary  0
         Mid-Career 90th Percentile Salary  37
         dtype: int64
```

```
In [32]: df.isnull().sum().sum()
```

```
Out[32]: 74
```

```
In [ ]: #from sklearn.impute import SimpleImputer
        #imp_mean = SimpleImputer(strategy='mean') #for median imputation replace 'mean' with 'median'
        #imp_mean.fit(df)
        #imputed_mean_df = imp_mean.transform(df)
        #Impute the values using scikit-learn SimpleImpute Class
```